

Towards Unstructured MAPF: Multi-Quadruped MAPF Demo

Rishi Veerapaneni*, Nikhil Sobanbabu*, Guanya Shi, Jiaoyang Li, Maxim Likhachev


Carnegie Mellon University
{rveerapa, nsobanba, guanyas, jiaoyanl, mlikhach}@andrew.cmu.edu

Abstract

Multi-Agent Path Finding (MAPF) in its most broad perspective focuses on finding collision free paths for general teams of agents in a shared environment. Theoretically, MAPF methods could solve a variety of multi-agent problems. However, MAPF research primarily focuses on simplified warehouse domains, i.e., gridworld with discrete spaces, discrete timesteps, and point-mass agents without kinematic constraints. Thus, the perception of MAPF is tied closely to gridworld and its assumptions, which limits its attractiveness to more broad domains. However, there are several ways to extend MAPF methods past these classical assumptions. To this end, our demo shows how MAPF techniques can be used to plan for a team of quadrupeds. Our system plans in continuous space, in continuous time, with realistic footprints, and incorporates dynamics constraints. See <https://youtu.be/ihVPEsN58t0> for additional details.

1 Introduction

Multi-Agent Path Finding (MAPF) focuses on finding collision-free paths for a team of agents in a shared workspace. The crown applications of MAPF are automated warehouse robotic systems which contain 100s-1000s of planar robots that need to transport items between different locations. Modern MAPF methods are extremely capable and can plan for 1000s of these agents in seconds. However, MAPF methods require many simplifying assumptions. In particular, most MAPF methods require a discretized world, discretized timesteps, point-mass agents, and no kinematic constraints. Thus, the perception of MAPF is that it is constrained to warehouses and not applicable to other complex robotics systems with fewer assumptions. To that end, our demo shows how to use MAPF for a team of quadrupeds. In particular, we plan in continuous space, continuous time, realistic footprints, and incorporate kinematic constraints for a team of quadrupeds. We note that we are not the first to relax these assumptions and that prior work, in particular dB-CBS (Moldagalieva et al. 2024), does so. Our demo seeks to emphasize these advancements with a real-world demo (Fig. 1) and simulated demo (Fig. 2) with multiple quadrupeds.



(a) 3 Quadrupeds



(b) 4 Quadrupeds

Figure 1: Real-world multi-quadruped experiments.

2 Related Works

Classic MAPF Formulation Multi-Agent Path Finding (MAPF) is the problem of finding collision-free paths for a group of N agents, that takes each agent i from its start location s_i^{start} to its goal location s_i^{goal} . In traditional 2D MAPF, the environment is discretized into grid cells, and time is broken down into discrete timesteps. Agents are allowed to move in any cardinal direction or wait in the same cell. A valid solution is a set of agent paths without vertex collisions (two agents at the same location at the same timestep) and edge collisions (two agents swapping locations). Thus, according to this construction, classical MAPF has the following restrictions: (1) Discretized locations, (2) Discretized timesteps, (3) Point-mass agents, and (4) No dynamics or kinematic constraints.

Variations on MAPF Assumptions There are several works that reduce the above assumptions. Large Agents MAPF (LA-MAPF) focuses on MAPF with agents with non-point-mass footprints (Li et al. 2019). LA-MAPF specifically introduces two different constraints in the context of Conflict-Based Search (Sharon et al. 2015). Several works have focused on MAPF with continuous time with non-point-mass footprints. Continuous time CBS (CCBS) (Andreychuk et al. 2019) plans for disk agents in continuous time on a 2^k connected grid. Discontinuity-Bounded CBS (dB-CBS) (Moldagalieva et al. 2024) uses CBS with a combination of heuristic search and optimization solvers to plan for a set of heterogeneous agents in continuous space, time, different footprints, and kinematic constraints.

Handling Execution Uncertainty After finding a collision-free MAPF solution, how do we execute these plans given that agents are not perfectly modeled and can have execution imperfections? In particular, there are two types of imperfections: (1) spatial uncertainty and (2)

*These authors contributed equally.




Figure 2: We visualize the pipeline for MAPF with quadrupeds (see Section 3).

temporal uncertainty. We deal with spatial uncertainty by assuming the agent’s controller has bounded tracking error and inflating the agent’s footprint during MAPF planning accordingly. We handle temporal uncertainty by using a Temporal Planning Graph (TPG) that encodes temporal dependencies between agents (Hönig et al. 2016).

3 MAPF for Multiple Quadrupeds

This section describes the technical details for planning and running multiple quadrupeds in a planar environment.

Modeling the Quadruped We note that many prior works that developed effective velocity tracking policies for legged robots. Thus given a sufficiently performant velocity controller, we can abstract away the quadrupeds 12 DoF leg joints (4 legs x 3 DoF per joints) but instead only represent the quadruped with a position (x, y) and a heading θ . The low-level RL policy enables the quadruped to track a 2D twist command consisting of linear translational velocity (v_x, v_y) , and angular velocity (w_z) . We incorporate heterogeneity in our framework with three different quadrupeds: Unitree Go2, Anybotics Anymal C, and Boston Dynamics’ Spot. However, we noticed that the pre-trained policies for the Unitree Go2 and Anymal C were poor in tracking high velocities and high acceleration trajectories. Hence, we trained custom policies for these two robots with extra regularization rewards and higher velocity command ranges.

MAPF Solver Given a set of start (x, y, θ) positions, goal (x, y, θ) positions, and rectangular footprints, the MAPF solver need to find collision-free (x, y, θ) paths while satisfying the agent’s holonomic kinematic properties. We use Conflict-Based Search (CBS) with a few modifications for practicality at the expense of optimality/bounded-suboptimality or completeness (Sharon et al. 2015). First, since we want to plan in continuous space and time, we use a Rapidly-Exploring Random Tree (RRT) single-agent planner (LaValle and Kuffner 2001). The RRT planner works in continuous space and time while avoiding obstacles, constraints, and incorporating the dynamic constraints. Second, since our agents have large footprints, we do not have edge conflicts but only deal with vertex conflicts. Since we work with continuous space and time, we detect collisions via a finely discretized collision checker. Third, given a detected collision between two rectangles at $(x_1^t, y_1^t, \theta_1^t)$ and $(x_2^t, y_2^t, \theta_2^t)$ respectively, we apply a simplified vertex constraint. We detect a collision point in the overlap between the agents and apply a constraint that each agent needs to

avoid overlapping with a spatial square containing the collision point for a time range (i.e., a space-time volume containing the collision point). We note that using this means that CBS is not complete. Finally, since we are mainly interested in finding feasible paths as opposed to finding minimal time paths, we use greedy-CBS which sorts the high-level queue via $(\# \text{num conflicts}, \text{sum of time})$ (Barer et al. 2014). Solutions took between 30-300 seconds to compute based on the number of quadrupeds and congestion.

Simulated Execution Given a collision-free set of paths, we execute them in Isaac Sim using our trained velocity tracking RL policies. The MAPF solution is first transformed into a TPG, which encodes the action dependencies across agents. Each quadruped in Isaac Sim follows this TPG, executing actions in sequence to reach the designated waypoints. To ensure smooth transitions between waypoints, the trajectory is refined using a Bézier spline. At each timestep, a proportional controller in the body frame computes the velocity commands for the RL policy based on the robot’s current position and the next intermediate waypoint within the spline. Figure 2 illustrates the planning and execution process of a four-agent system navigating from their start positions to the corresponding goal locations without any collisions. We further evaluated our pipeline on maps with 5-, 8-, and 12-agent systems. The average execution times observed were 113.32 ± 4.79 , 198.60 ± 12.44 , and 312.86 ± 23.76 seconds, respectively. These results were obtained over three independent trials for each configuration, with each trial using randomized, heterogeneous agents.

Real-World Execution We also verified this system on teams of 2-4 Unitree Go1s in a real-world physical set-up. We followed the same pipeline except used an A* single-agent planner as it produced paths easier to follow than the RRT, and during execution replaced the Bézier spline refinement with linear interpolation. Obstacles were hand defined in the 2D occupancy map to match the real-world set-up, and the quadrupeds were tracked using OptiTrack motion capture system. Incorporating a TPG was crucial, as differences in gait and tracking accuracy across quadrupeds often caused deviations from the nominal path, yet the TPG enabled dynamic pausing without requiring online replanning. Most failures stemmed from hardware issues (e.g., leg malfunctions) rather than MAPF execution. Future work could focus on dynamically handling such failures or leveraging the quadrupeds’ full mobility.

Acknowledgements

The authors thank Shiqi Liu, Yaru Niu, and Professor Ding Zhao for their critical help in conducting the real-world quadruped demonstrations, and Yue Zhang for her critical help in presenting our work at ICAPS.

References

- Andreychuk, A.; Yakovlev, K.; Atzmon, D.; and Stern, R. 2019. Multi-Agent Pathfinding with Continuous Time. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI-19*, 39–45.
- Barer, M.; Sharon, G.; Stern, R.; and Felner, A. 2014. Sub-optimal variants of the conflict-based search algorithm for the multi-agent pathfinding problem. In *Seventh Annual Symposium on Combinatorial Search*.
- Hönig, W.; Kumar, T. S.; Cohen, L.; Ma, H.; Xu, H.; Ayanian, N.; and Koenig, S. 2016. Multi-agent path finding with kinematic constraints. In *Twenty-Sixth International Conference on Automated Planning and Scheduling*.
- LaValle, S. M.; and Kuffner, J. J. 2001. Rapidly-exploring random trees: Progress and prospects. *Algorithmic and computational robotics*, 303–307.
- Li, J.; Surynek, P.; Felner, A.; Ma, H.; Kumar, T. K. S.; and Koenig, S. 2019. Multi-Agent Path Finding for Large Agents. *Proceedings of the AAAI Conference on Artificial Intelligence*, 33(01): 7627–7634.
- Moldagalieva, A.; Ortiz-Haro, J.; Toussaint, M.; and Hönig, W. 2024. db-CBS: Discontinuity-bounded conflict-based search for multi-robot kinodynamic motion planning. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, 14569–14575. IEEE.
- Sharon, G.; Stern, R.; Felner, A.; and Sturtevant, N. R. 2015. Conflict-based search for optimal multi-agent pathfinding. *Artificial Intelligence*, 219: 40–66.